

```
In [1]: import pandas as pd
        from sklearn.model_selection import train_test_split
        from sklearn.feature_extraction.text import TfidfVectorizer
        from sklearn.naive_bayes import BernoulliNB
        from sklearn.linear_model import LogisticRegression
        from sklearn.metrics import accuracy_score
        from sklearn.neural_network import MLPClassifier
```

1

```
In [2]: data = pd.read_csv("federalist.csv")
        data['author'] = data.author.astype('category')
        data.head()
```

```
Out[2]:
```

	author	text
0	HAMILTON	FEDERALIST. No. 1 General Introduction For the...
1	JAY	FEDERALIST No. 2 Concerning Dangers from Forei...
2	JAY	FEDERALIST No. 3 The Same Subject Continued (C...
3	JAY	FEDERALIST No. 4 The Same Subject Continued (C...
4	JAY	FEDERALIST No. 5 The Same Subject Continued (C...

```
In [3]: data.author.value_counts()
```

```
Out[3]:
```

HAMILTON	49
MADISON	15
HAMILTON OR MADISON	11
JAY	5
HAMILTON AND MADISON	3

Name: author, dtype: int64

```
In [4]: data.dtypes
```

```
Out[4]:
```

author	category
text	object

dtype: object

2

```
In [5]: X = data.text
        y = data.author

        X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, train_size=0.8)
        X_train.shape, X_test.shape, y_train.shape, y_test.shape
```

```
Out[5]: ((66,), (17,), (66,), (17,))
```

3

In [6]:

```
vectorizer = TfidfVectorizer(stop_words='english')

# vectorize
X_train = vectorizer.fit_transform(X_train) # fit the training data
X_test = vectorizer.transform(X_test) # transform only

print('train size:', X_train.shape)
print(X_train.toarray()[:5])

print('\ntest size:', X_test.shape)
print(X_test.toarray()[:5])
```

train size: (66, 7727)

[0.	0.	0.03056353	...	0.	0.	0.]
[0.	0.	0.	...	0.	0.	0.]
[0.	0.	0.	...	0.	0.	0.]
[0.	0.	0.	...	0.	0.	0.]
[0.	0.	0.	...	0.03999681	0.	0.]]

test size: (17, 7727)

[0.	0.	0.	...	0.	0.	0.]
[0.	0.	0.	...	0.02406012	0.	0.]
[0.	0.	0.	...	0.	0.	0.]
[0.	0.	0.	...	0.	0.	0.]
[0.	0.	0.	...	0.	0.	0.]]

4

In [7]:

```
naive_bayes = BernoulliNB()
naive_bayes.fit(X_train, y_train)
```

Out[7]: BernoulliNB()

In [8]:

```
pred = naive_bayes.predict(X_test)
print('accuracy score: ', accuracy_score(y_test, pred))
```

accuracy score: 0.5882352941176471

5

In [9]:

```
X = data.text
y = data.author
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, train_size=0.8)
X_train.shape, X_test.shape, y_train.shape, y_test.shape
```

Out[9]: ((66,), (17,), (66,), (17,))

In [10]:

```
vectorizer = TfidfVectorizer(stop_words='english', max_features=1000, ngram_range=(1,2))
```

```
# vectorize
X_train = vectorizer.fit_transform(X_train) # fit the training data
X_test = vectorizer.transform(X_test) # transform only

print('train size:', X_train.shape)
print(X_train.toarray()[:5])

print('\ntest size:', X_test.shape)
print(X_test.toarray()[:5])
```

```
train size: (66, 1000)
[[0.02001784 0.02001784 0.01929126 ... 0.03268402 0.01766216 0.          ]
 [0.02282457 0.02282457 0.          ... 0.03726669 0.02013859 0.          ]
 [0.01136023 0.01136023 0.          ... 0.10510731 0.01002337 0.          ]
 [0.          0.          0.          ... 0.01044773 0.          0.          ]
 [0.01676058 0.01676058 0.08076111 ... 0.01824383 0.01478821 0.04437088]]

test size: (17, 1000)
[[0.02956605 0.02956605 0.0569858 ... 0.03218254 0.02608674 0.          ]
 [0.0233583  0.0233583  0.02251047 ... 0.02542542 0.02060951 0.03091864]
 [0.02531887 0.02531887 0.02439988 ... 0.02755951 0.02233937 0.          ]
 [0.0241714  0.0241714  0.          ... 0.02631049 0.02132693 0.          ]
 [0.01918632 0.01918632 0.01848992 ... 0.02088424 0.01692849 0.          ]]
```

```
In [11]: naive_bayes = BernoulliNB()
naive_bayes.fit(X_train, y_train)
```

```
Out[11]: BernoulliNB()
```

```
In [12]: pred = naive_bayes.predict(X_test)
print('accuracy score: ', accuracy_score(y_test, pred))
```

```
accuracy score:  0.9411764705882353
```

6

```
In [13]: classifier = LogisticRegression()
classifier.fit(X_train, y_train)

# evaluate
pred = classifier.predict(X_test)
print('accuracy score: ', accuracy_score(y_test, pred))
```

```
accuracy score:  0.5882352941176471
```

```
In [14]: classifier = LogisticRegression(solver='lbfgs', class_weight='balanced')
classifier.fit(X_train, y_train)

# evaluate
pred = classifier.predict(X_test)
print('accuracy score: ', accuracy_score(y_test, pred))
```

```
accuracy score:  0.7058823529411765
```

7

In [15]:

```
classifier = MLPClassifier(solver='lbfgs', alpha=1e-5,  
                           hidden_layer_sizes=(100,), random_state=1)  
classifier.fit(X_train, y_train)  
  
pred = classifier.predict(X_test)  
print('accuracy score: ', accuracy_score(y_test, pred))
```

accuracy score: 0.7058823529411765

In [16]:

```
classifier = MLPClassifier(solver='lbfgs', alpha=1e-5,  
                           hidden_layer_sizes=(50,), random_state=1)  
classifier.fit(X_train, y_train)  
  
pred = classifier.predict(X_test)  
print('accuracy score: ', accuracy_score(y_test, pred))
```

accuracy score: 0.7647058823529411