

Accuracy:

The output of our model's accuracy was:

```
Accuracy: 97.33333333333334
```

For a probabilistic language model, this is a very good accuracy score. One of the reasons such a high score was received might be due to the fact that the training and testing data are very similar in nature. Both are composed of political discussions, so the model has a relatively narrow subset of vocabulary to learn from, which is much easier than learning the entirety of a language.

Incorrect Classifications:

As for the incorrectly classified items, they they were outputted as:

```
Incorrectly classified: [24, 44, 92, 187, 191, 247, 277, 279]
```

where line 1 is the first line of the document. Looking at each of these lines, there aren't many similar characteristics between them, so we can assume they were each misclassified for different reasons. For example, line 92 contains only three words, so its short length may be the cause of its misclassification.